

Lecture #2: Branch Prediction

Parallel Computer Systems

Lieven Eeckhout

Academiejaar 2014-2015
Ghent University

Instruction fetch

- Goal is to fetch as many insns per cycle as possible
- Branch prediction is super important to fill the pipeline with correct-path instructions (useful work)
- Importance increases with deeper pipelines

Branch prediction

- Goal is to predict the branch direction and target address, and start fetching and executing insns along the predicted path
- Key observation:
 - Branches exhibit temporal locality
 - Predicting branch behavior
 - Keep track of past history
 - Predict the future based on the past
 - Branch behavior is predictable
 - typically over 90%, 95% or 99% of all dynamically executed branches are correctly predicted

Overview

- Branch direction prediction
 - Static
 - Dynamic
- Branch target prediction
- How to handle mispredictions?

Branch Direction Prediction

- Applies to conditional branches
- Static prediction
 - Static = before program execution
 - One prediction (taken/not-taken) per static branch in the program binary
 - Via software: compiler or programmer
- Dynamic prediction
 - Dynamic = during program execution
 - Multiple predictions per static branch, depending on history (= outcomes of prior branch executions) of that particular branch or even other branches
 - Done in hardware

Branch address vs. branch target address

```
000000012002e530 <__start>:
12002e530: f0 ff de 23    lda    sp,-16(sp)
12002e534: 08 00 fe b7    stq    zero,8(sp)
12002e538: 00 00 20 c0    br     t0,12002e53c <__start+0xc>
12002e53c: 10 00 1e a2    ld1    a0,16(sp)
12002e540: 18 00 3e 22    lda    a1,24(sp)
12002e544: 00 20 a1 27    ldah   gp,8192(t0)
12002e548: 52 06 11 42    s8addq a0,a1,a2
12002e54c: 12 14 41 42    addq   a2,0x8,a2
12002e550: 24 28 bd 23    lda    gp,10276(gp)
12002e554: 13 04 52 46    mov    a2,a3
12002e558: 00 00 33 a4    ldq    t0,0(a3)
12002e55c: 13 14 61 42    addq   a3,0x8,a3
12002e560: fd ff 3f f4    bne    t0,12002e558 <__start+0x28>
```

under Linux/UNIX/OS X: `objdump -d a.out`

Static Branch Prediction

- Advantages
 - Easy to implement
 - Little HW is needed
- Disadvantages
 - Provides the same prediction regardless of input and/or dynamic execution behavior
- Three flavors
 - rule-based
 - program-based
 - profile-based

Rule-based

- Always not-taken
 - Simple HW: sequential fetch
- Always taken
 - HW is more complex because of unknown branch target
 - Branch target is known at the decode stage
 - May lead to a *bubble* ('lost cycle') in the pipeline
- BTFNT
 - *Backward taken - forward not-taken*
 - Branches to smaller addresses are typically backward loop branches and are typically taken

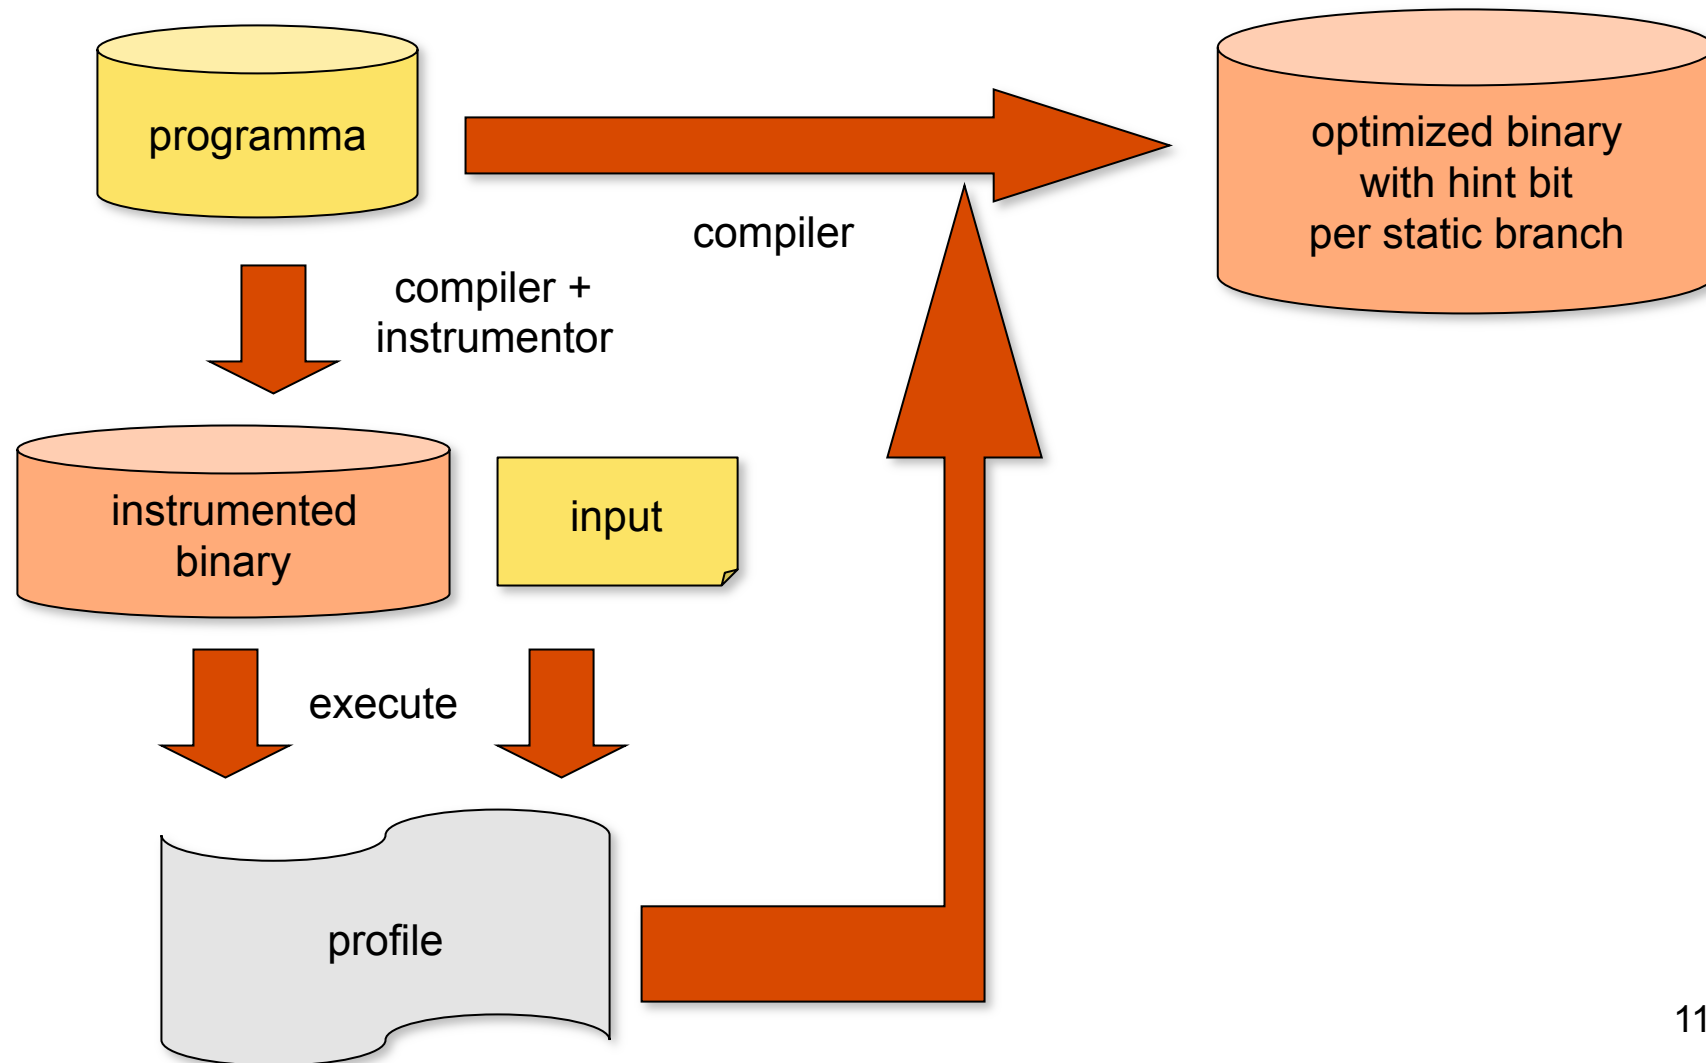
Program-based

- Ball en Larus heuristics (1993)
 - Requires a hint bit in instruction opcode
 - Branch direction is estimated based on program structure
 - Examples:
 - Predict loop branches to be taken
 - When comparing a pointer to NULL, predict branch direction to non-NULL path
 - When comparing two pointers, predict branch direction to path representing pointer inequality
 - Typically more accurate than rule-based

Profile-based

- Execute instrumented binary with a given training input to collect profile information
 - Count how often a static branch is taken/not-taken
- Use profile information during recompilation and add hint bits
 - Predict taken if branch has higher probability than 50% to be taken; predict not-taken otherwise
- Typically more accurate than rule- and program based

Profile-based



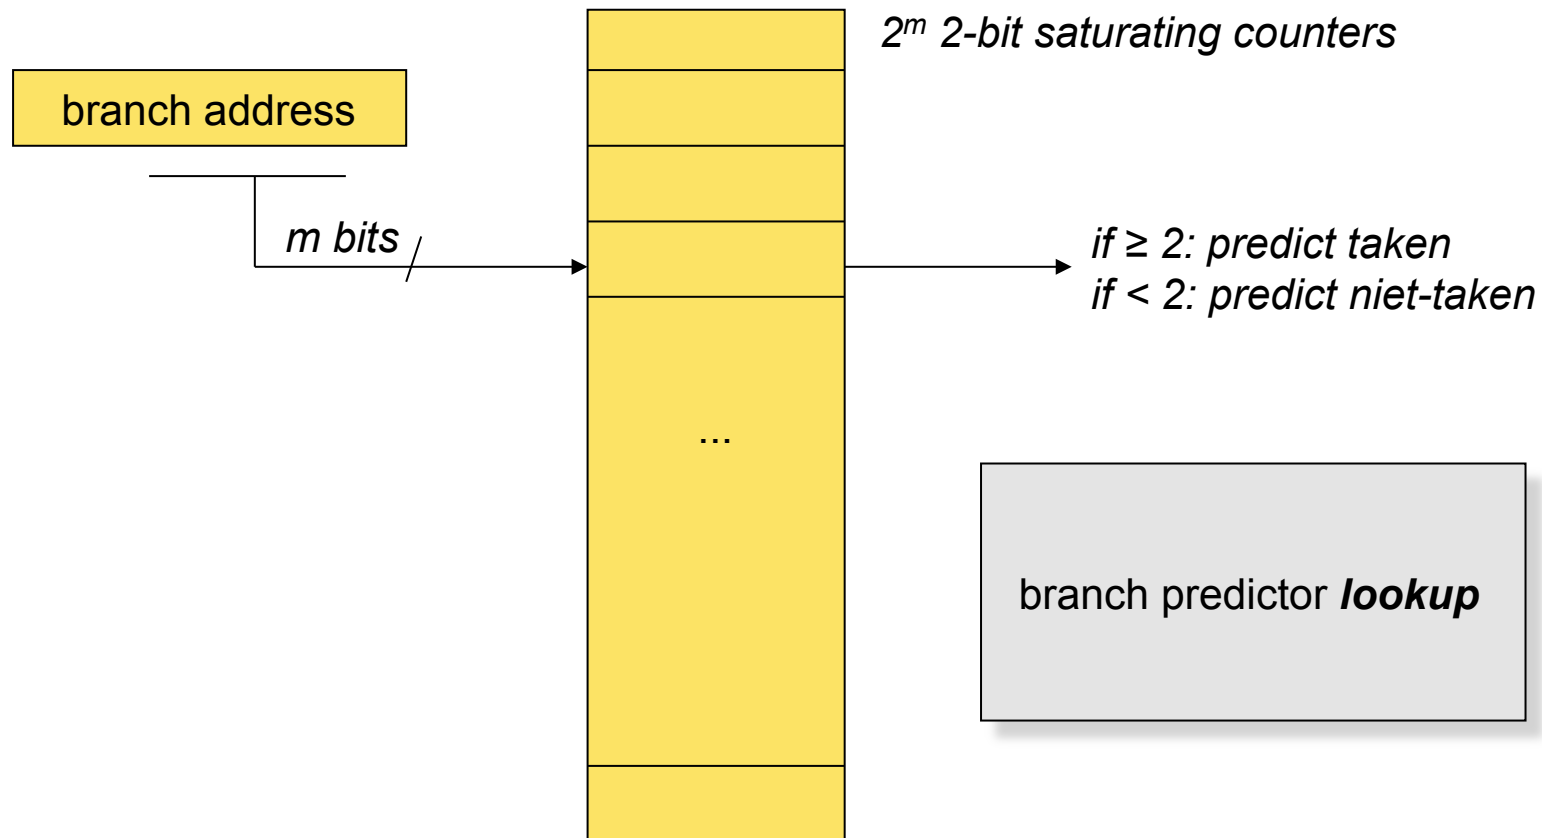
Overview

- Branch direction prediction
 - Static
 - Dynamic
- Branch target prediction
- How to handle mispredictions?

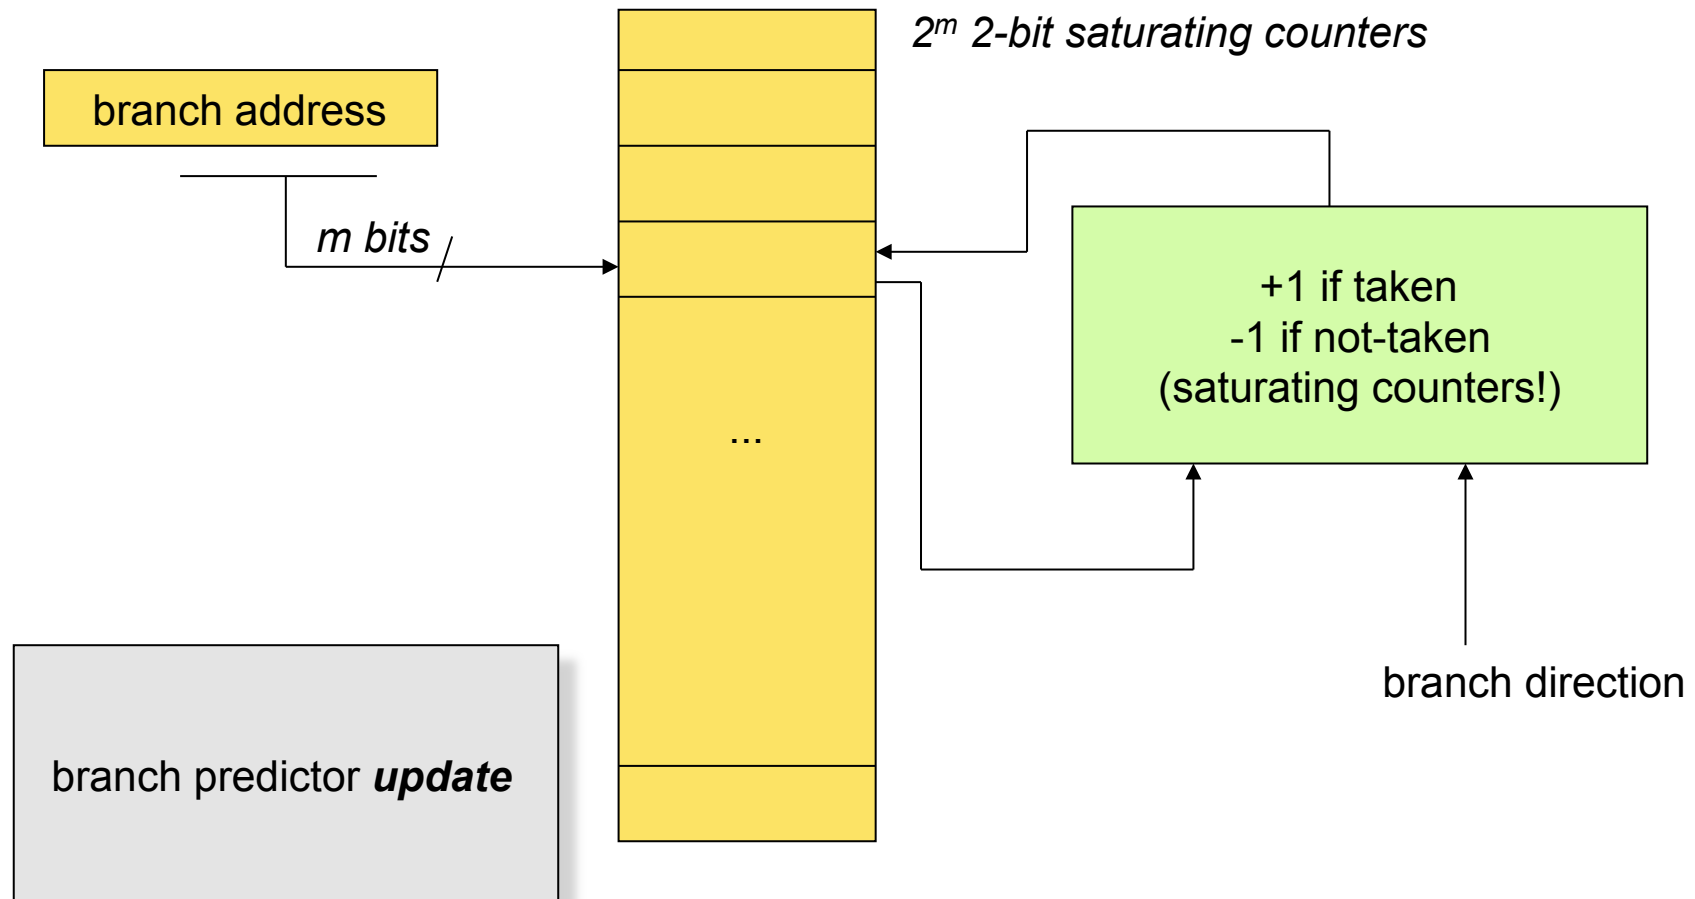
Dynamic branch prediction

- More accurate than static branch prediction
 - 80%-97% (dynamic) vs. 50%-80% (static)
- Some branches are hard to predict statically, but are easily predicted dynamically
 - Some examples
 - Not-taken during first half of execution, and taken during second half
 - Alternating taken/not-taken
- Takes into account branch context!
 - the branch's own history (*local history*)
 - other branches' histories (*global history*)

Bimodal predictor



Bimodal predictor



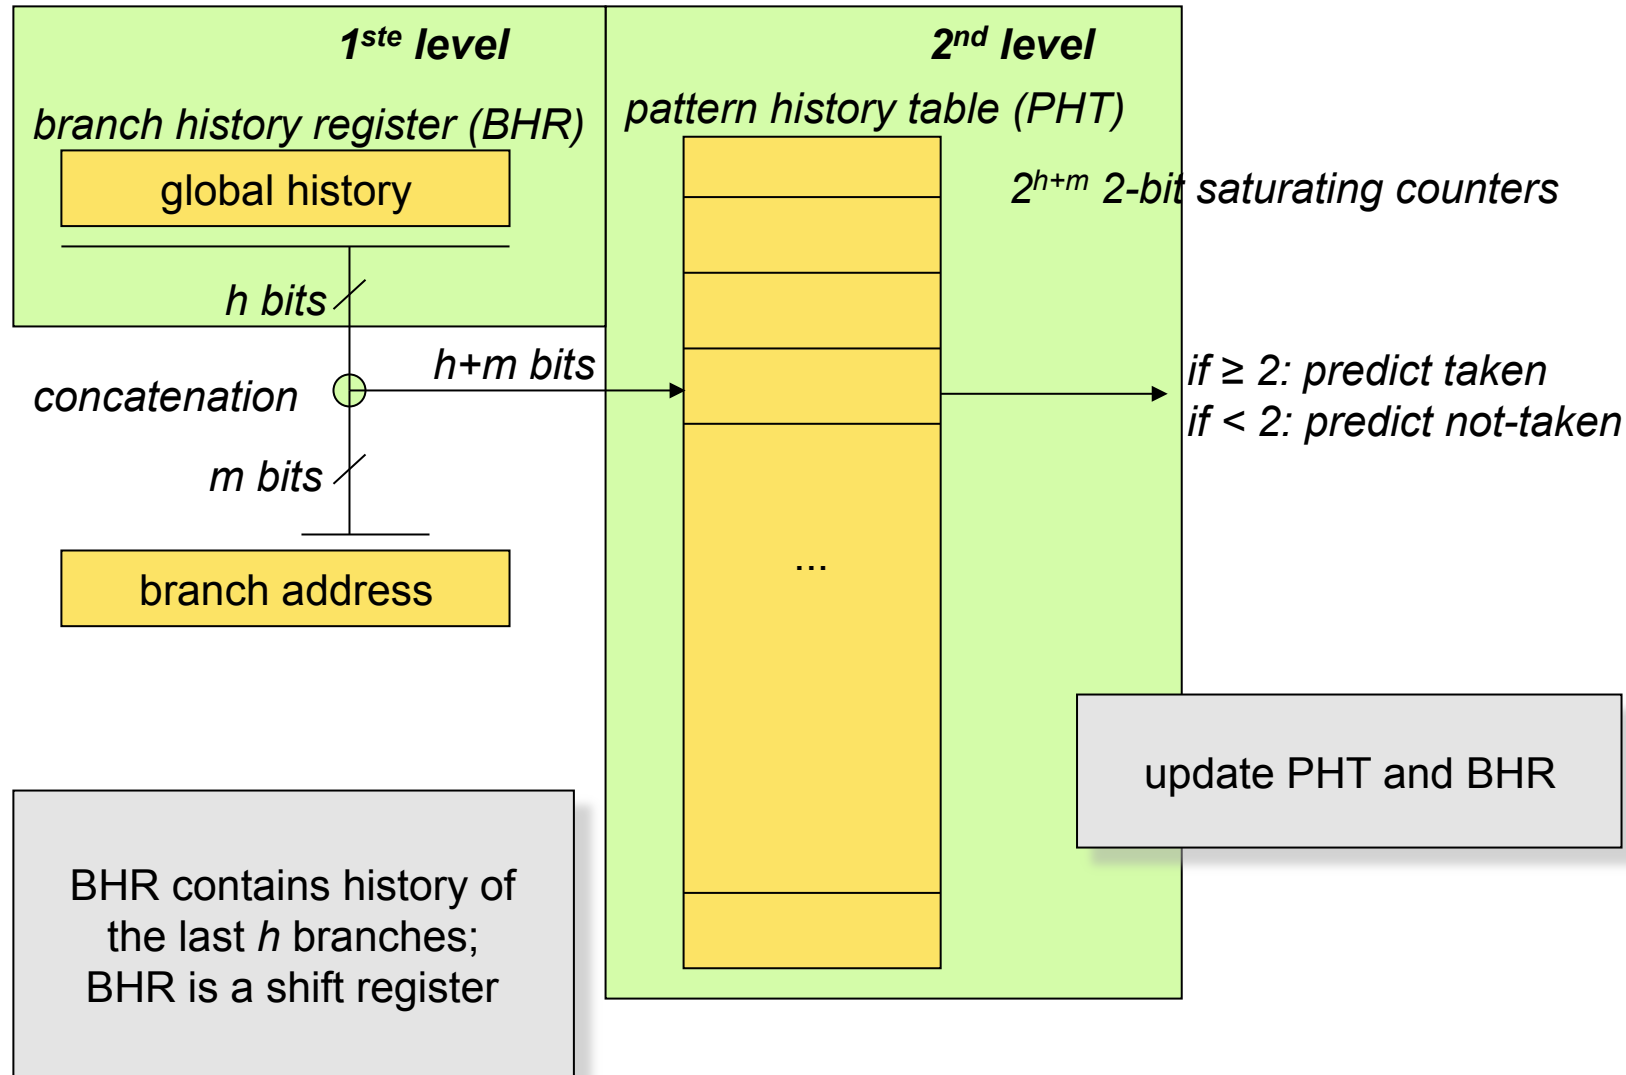
Example bimodal predictor

<i>old state</i>	<i>prediction</i>	<i>branch direction</i>	<i>new state</i>
00	0	0	00
00	0	0	00
00	0	1	01
01	0	1	10
10	1	1	11
11	1	1	11
11	1	0	10
10	1	1	11
11	1	1	11
11	1	0	10
10	1	0	01

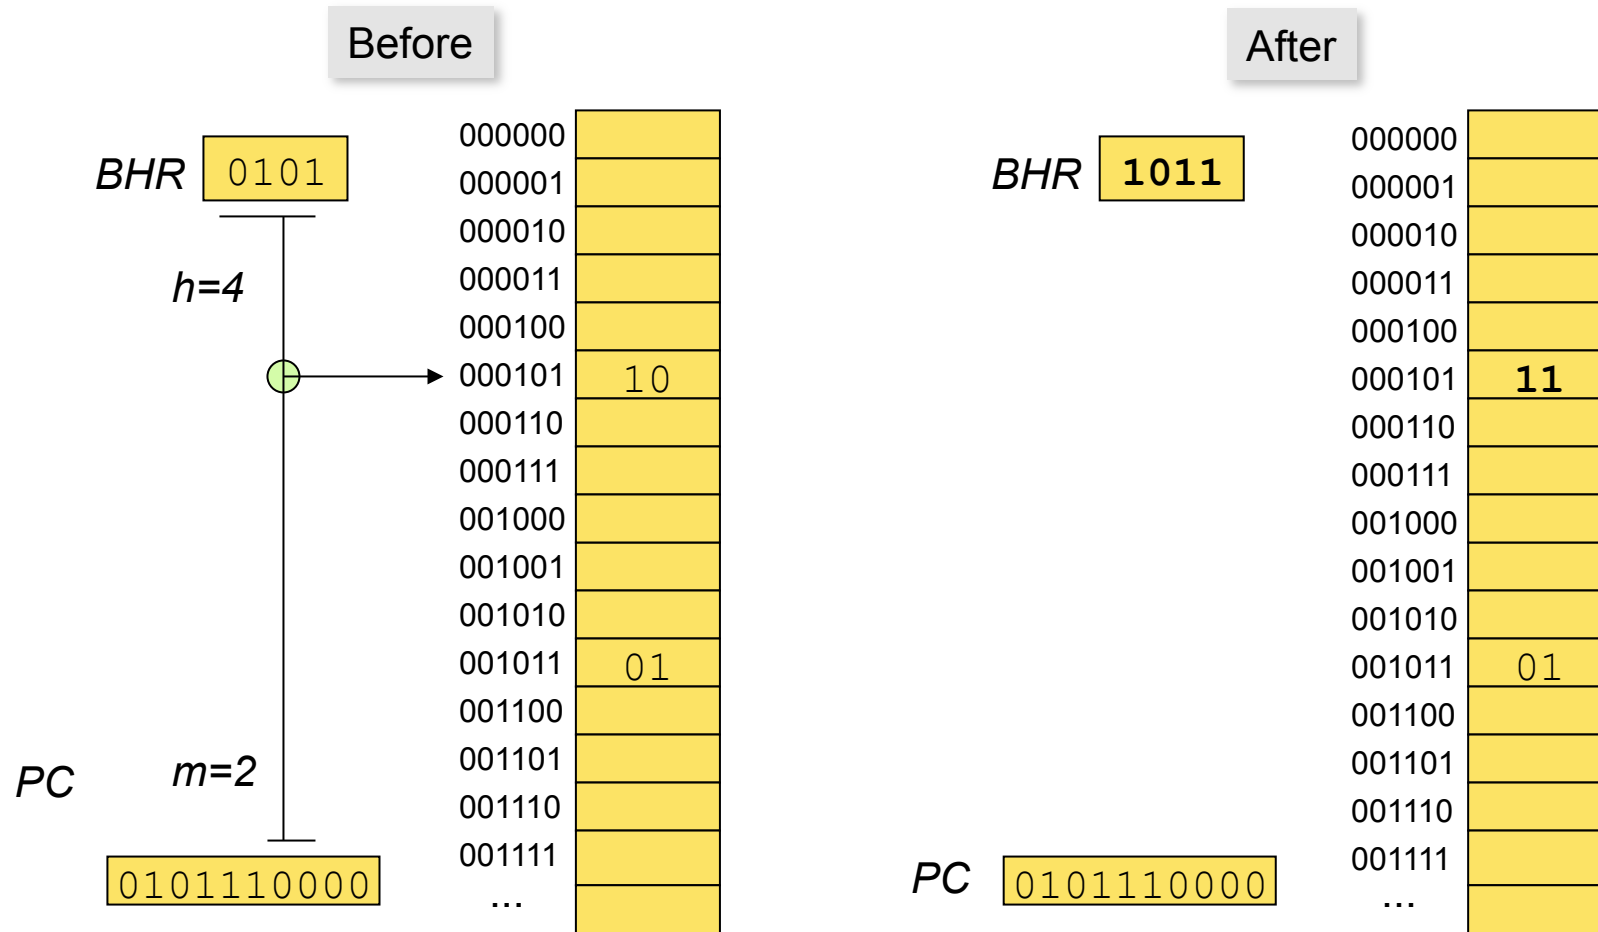
Two-level predictors

- *Correlating branch predictor*
- Uses (two levels of) branch history to make a prediction
 - Bimodal predictor uses the branch address only
 - Two-level predictors also use
 - *Local history* of that particular branch, or
 - *Global history* of all prior branches
 - history = outcome of prior branch executions
 - history acts as a shift register: newest branch outcome is inserted; oldest branch outcome is discarded
- Proposed by Tse-Yu Yeh and Yale N. Patt in 1991

... using global history



Example



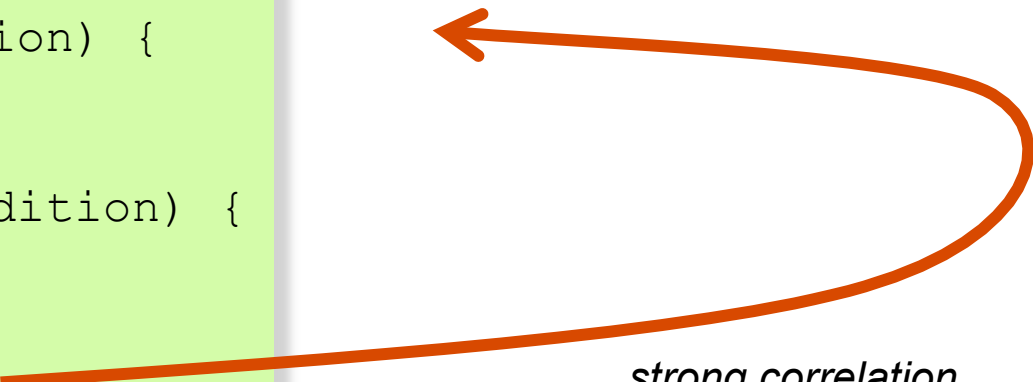
Assume: branch direction is taken

Why does this work?

- Correlation among branches
 - For example, if a particular branch is taken (not-taken), probability may be high that next branches are also taken (not-taken), or vice versa
- Concrete examples
 - Branch conditions that depend on the same variable
 - A variable is control dependent on branch; a subsequent branch is data dependent on the variable; both branches will be correlated

Example

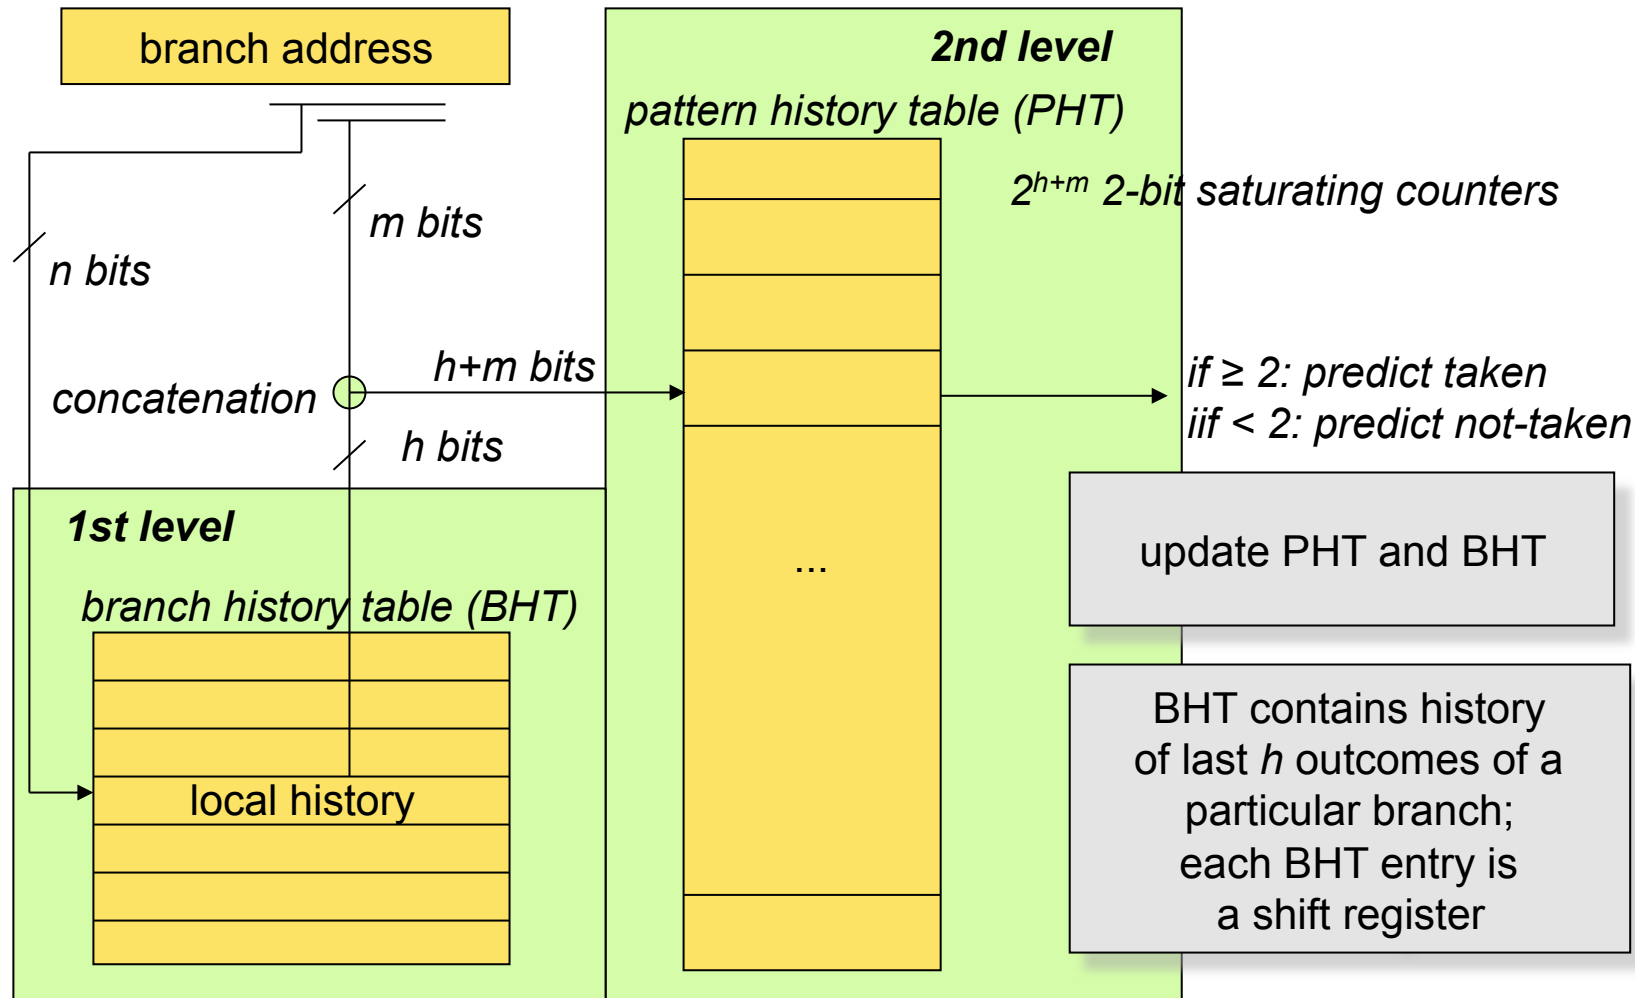
```
x = 0;  
if (some_condition) {  
    x = 3;  
}  
if (another_condition) {  
    y += 19;  
}  
if (x <= 0) {  
    do_something ();  
}
```



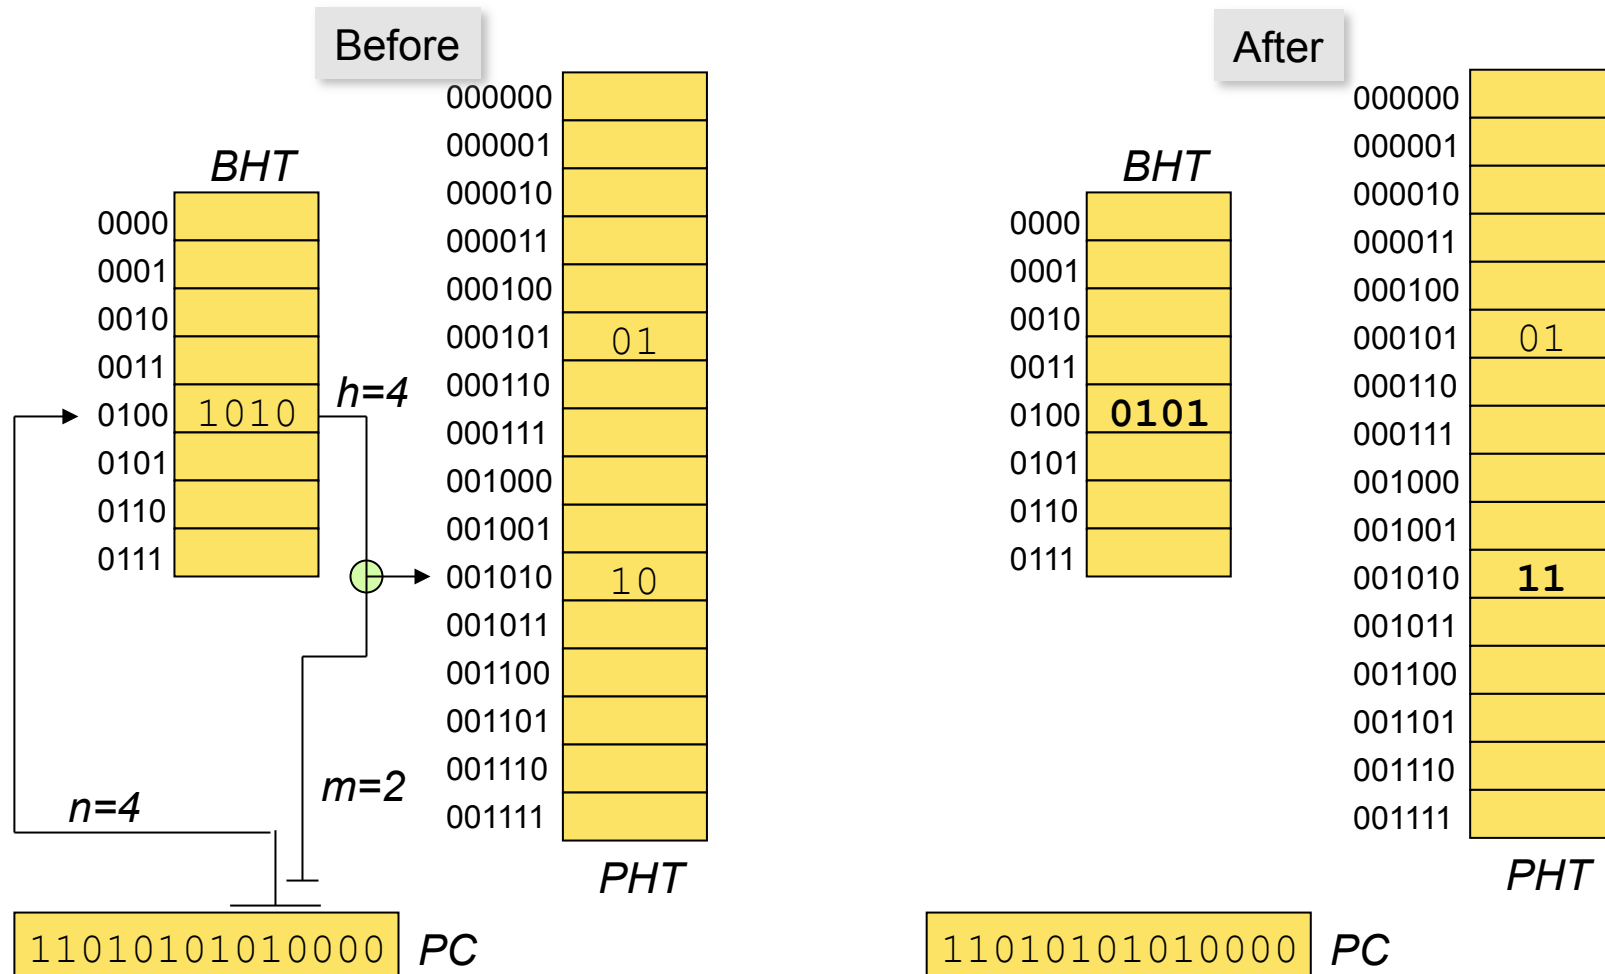
*strong correlation
between 1st and 3rd branch*

Note: there is no correlation between
the 2nd and 3rd branch → branch
prediction needs to learn!

... using local history



Example



Assume: branch direction is taken

Why does this work?

- Local behavior of a branch
- Examples
 - Alternating: 01010101...
 - Loops with limited number of iterations:
111011101110...

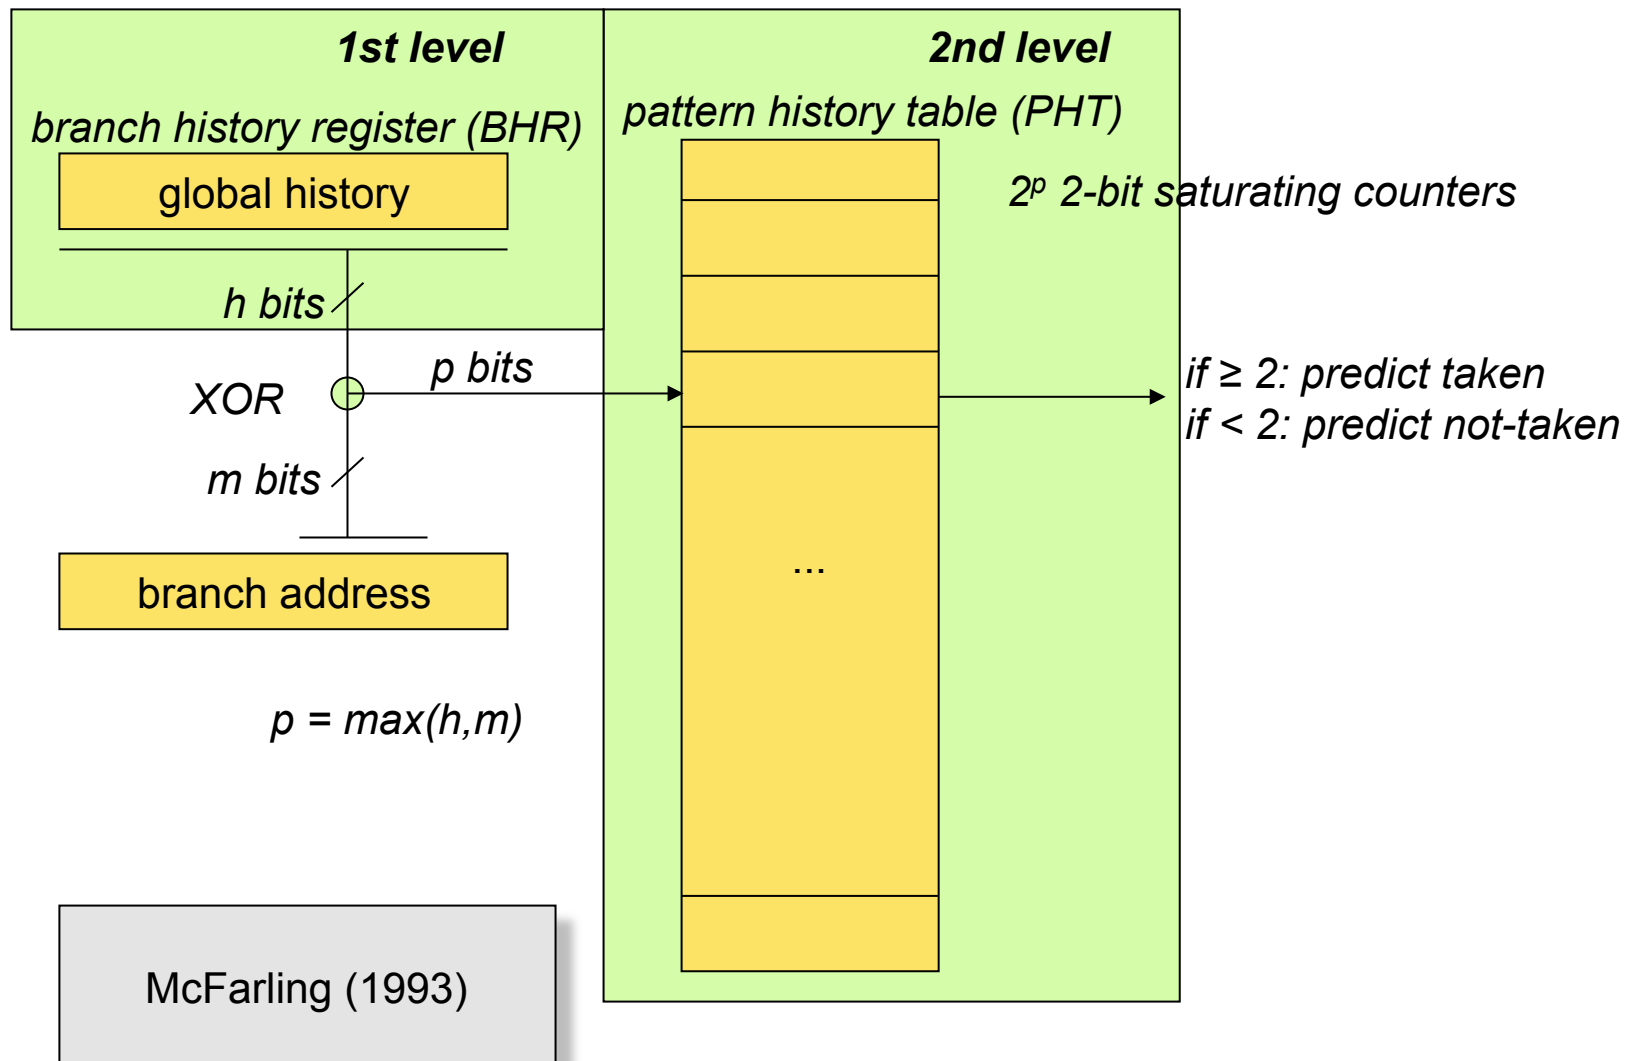
Two-level predictor implementations

- If $m=0$
 - Global pattern history table (*gPHT*)
 - All branches share the same PHT
- If $m \neq 0$
 - Per-address pattern history table (*pPHT*)
 - PHT is partitioned based on branch address bits
- Four variations: GAg, GAp, PAg, PAp
 - ‘G’ global history, ‘P’ per-address (local)
 - ‘A’ adaptive
 - ‘g’ \rightarrow gPHT and ‘p’ \rightarrow pPHT

Good configurations

- Obtained through experimental evaluation by Yeh/Patt (1991)
- GAg
 - BHR: 18 bits
 - PHT: $2^{18} \times 2$ bits
- PAg
 - BHT: $2^{11} \times 12$ bits
 - PHT: $2^{12} \times 2$ bits
- PAp
 - BHT: $2^{11} \times 6$ bits
 - PHT: $2^9 \times 2^6 \times 2$ bits
- Prediction accuracy of 97%

gshare



GAp vs. gshare

- GAp needs to make a choice
 - How many bits to concatenate from BHR vs. branch address?
 - e.g., 5 bits from BHR and 5 bits from branch address
- gshare doesn't need to make this choice
 - e.g., 10 bits from BHR and 10 bits from branch address
 - More information is used to index the PHT

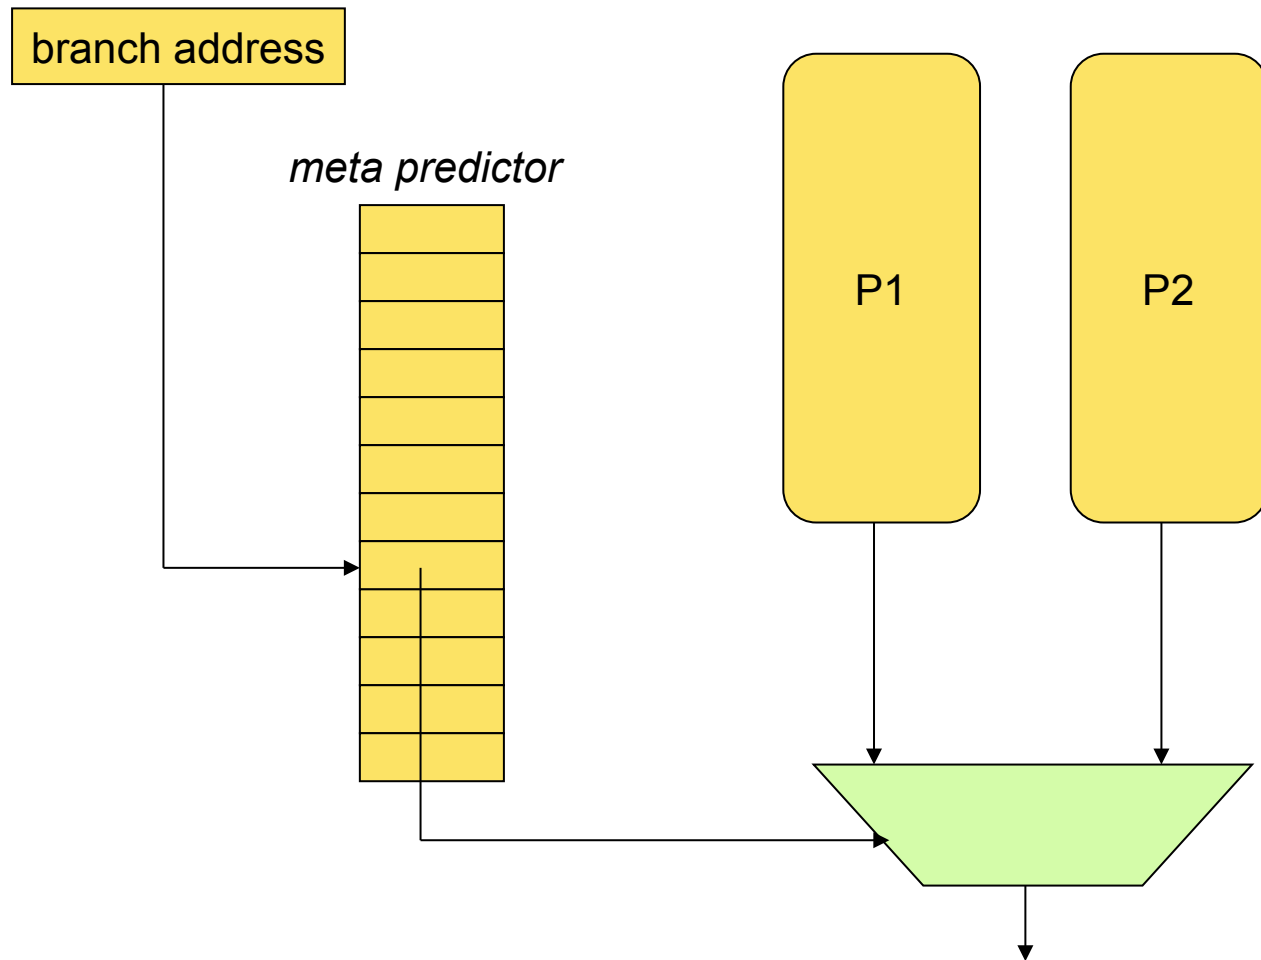
Towards hybrid branch prediction

- Reasons for mispredictions?
- Branches may be hard to predict
 - During training phase of the branch predictor, e.g., after a context switch
- Interference or aliasing
 - Mostly negative interference
 - Branch predictor is of too limited size for the number of branches
 - Hence, branches with different behavior may be mapped onto the same PHT entry
- Branch behavior doesn't match branch predictor type

Hybrid branch prediction

- Idea: Combine different types of branch predictors and learn which one is most accurate for which branch

Tournament predictor



Tournament predictor

- Meta predictor determines which predictor will be followed
 - If $< 2 \rightarrow P1$; If $\geq 2 \rightarrow P2$
- Update meta predictor
 - If both correct or incorrect: do nothing
 - If P1 is correct and P2 incorrect: decrement
 - If P1 is incorrect and P2 is correct: increment
- Both predictors are updated
 - irrespective of whether P1 or P2 was followed

Tournament predictor (cont.)

- Various flavors exist
- Any predictor can be part of a hybrid predictor
- Indexing of meta predictor is to be chosen
- Hybrid predictor typically combines local and global history based predictors

Alpha 21264

- Hybrid predictor
 - PAg
 - 1st level: 1K 10-bit entries
 - 2nd level: 1K 3-bit saturating counters
 - GAg
 - 4K 2-bit saturating counters
 - 12-bit global history
 - Meta predictor
 - 4K 2-bit saturating counters
 - Indexed using 12-bit global history

IBM POWER4

- Hybrid predictor
 - Bimodal predictor
 - 16K 1-bit saturating counters
 - Gshare
 - 16K 1-bit saturating counters
 - 11-bit global history
 - Meta predictor
 - 16K 1-bit saturating counters
 - gshare indexing

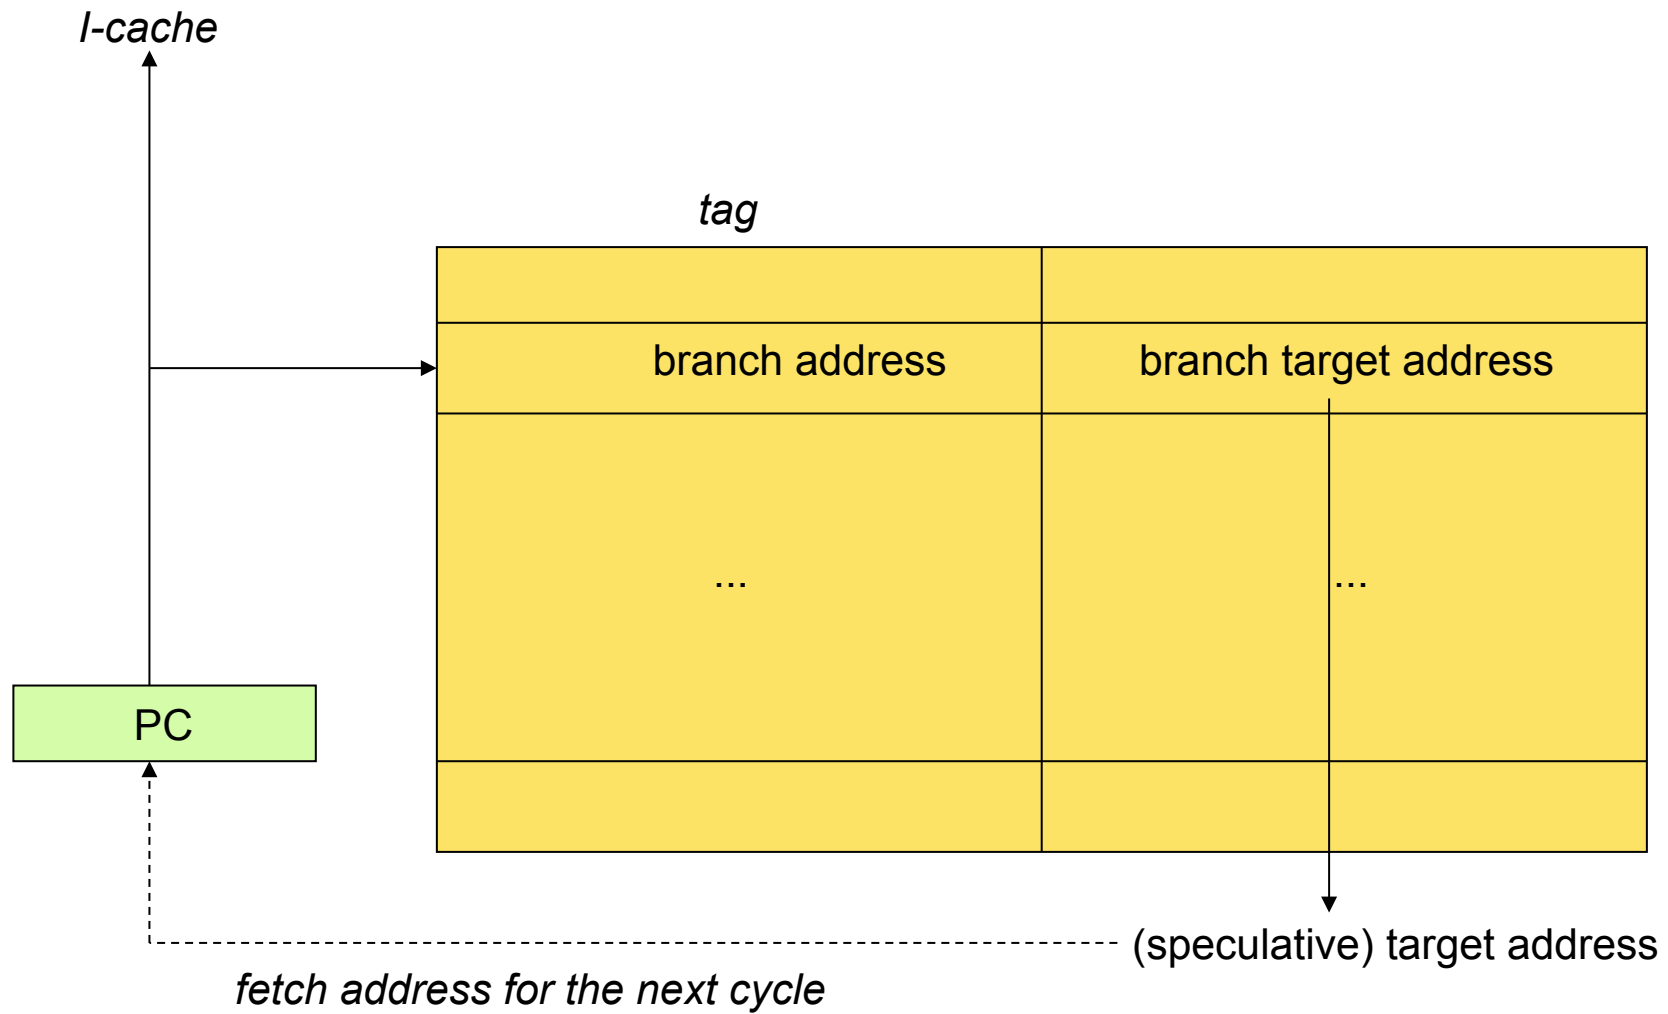
Overview

- Branch direction prediction
 - Static
 - Dynamic
- Branch target prediction
- How to handle mispredictions?

Branch target prediction

- Branch target buffer (BTB) or Branch target address cache (BTAC)
- Small (set-associative) cache
- Idea:
 - Input = branch address
 - Output = branch target address
 - Accessed in the same cycle as the I-cache access
 - Branch target address is the fetch address in the next cycle
- Done for both conditional as well as non-conditional branches


Branch Target Buffer



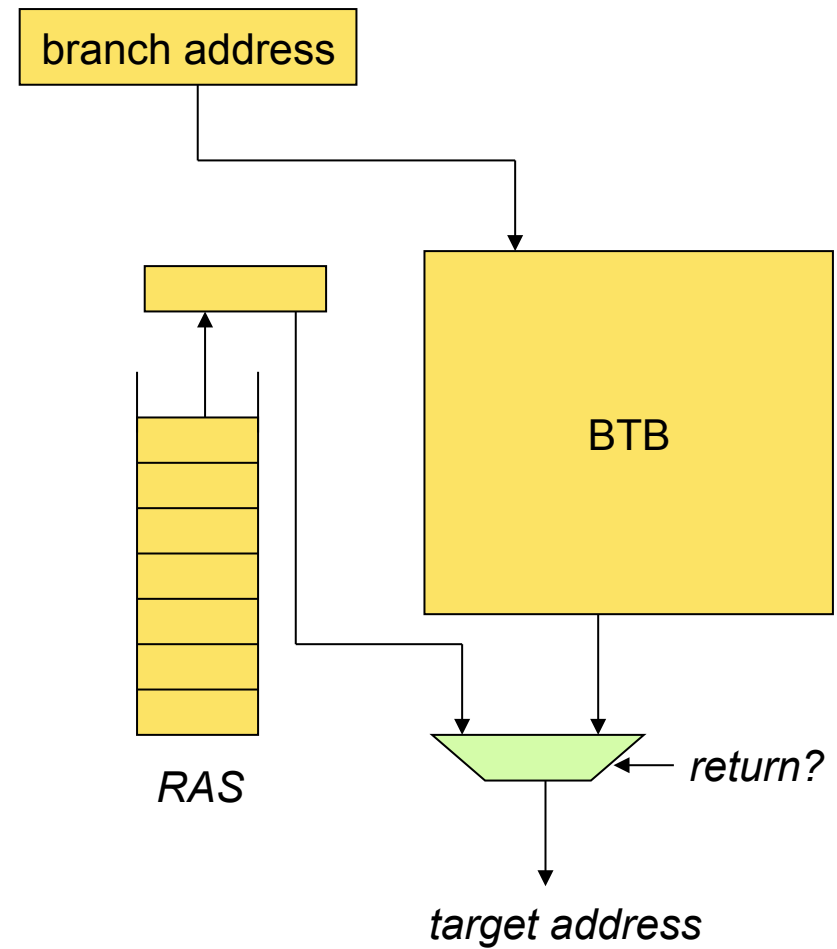
Return Address Stack

- Function calls are easy to predict
- Function returns are much more complicated to predict
- Solution: Return Address Stack (RAS)
 - by Kaeli and Emma in 1991
 - In Pentium 4: 16 elementen
 - Circular buffer
 - Push function return address on RAS; pop the address upon a return
 - If depth of function call stack > RAS size: incorrect predictions

Return Address Stack

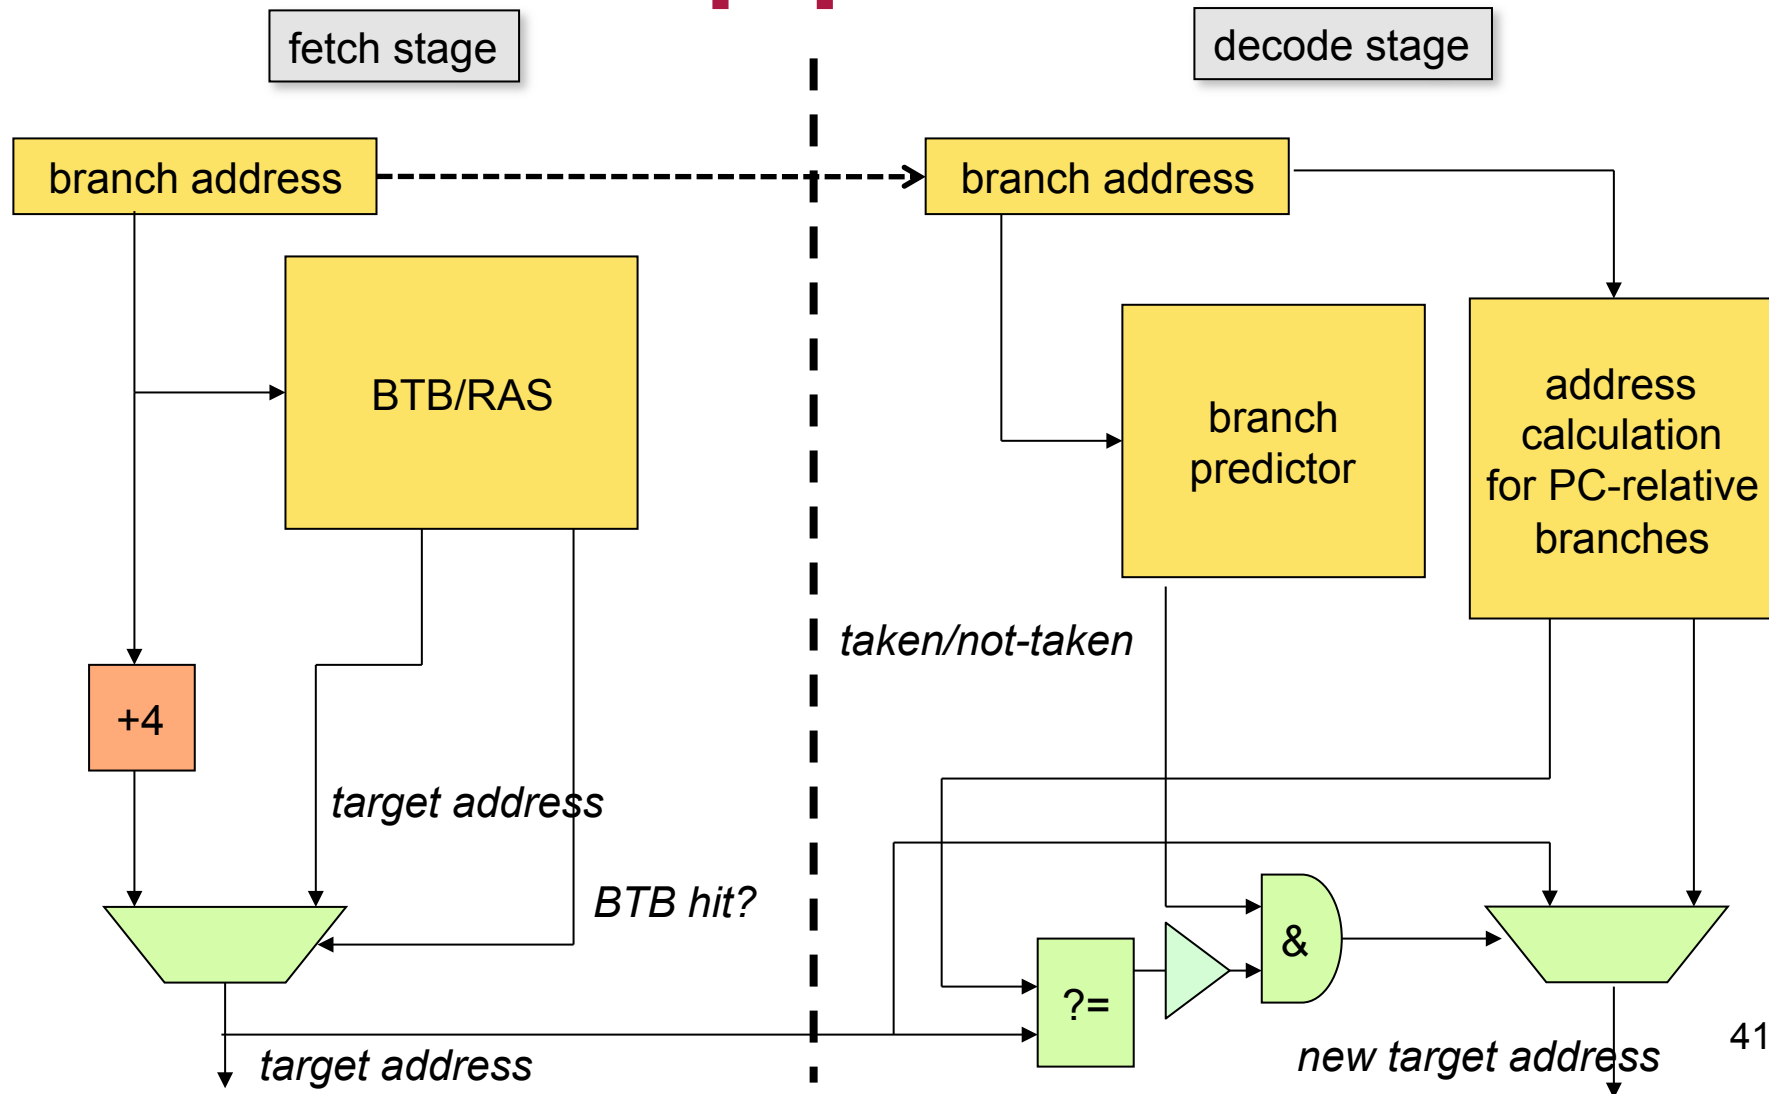


In case of a function call

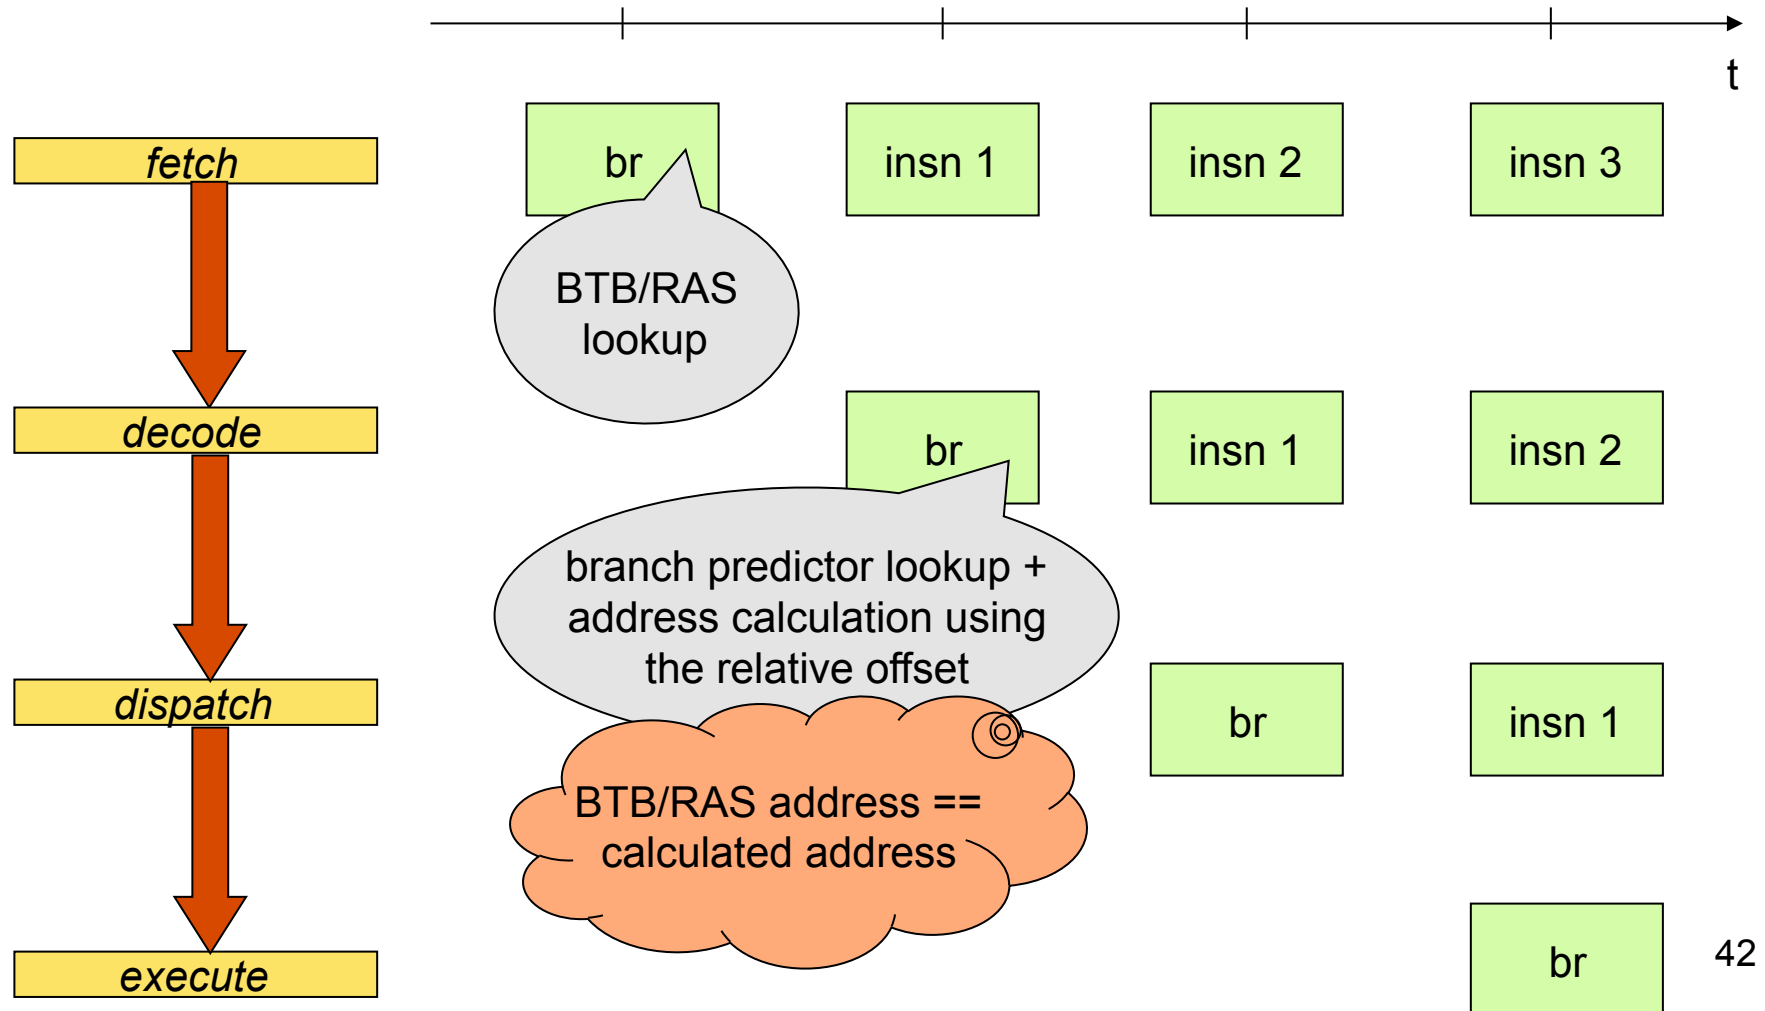


In case of a return

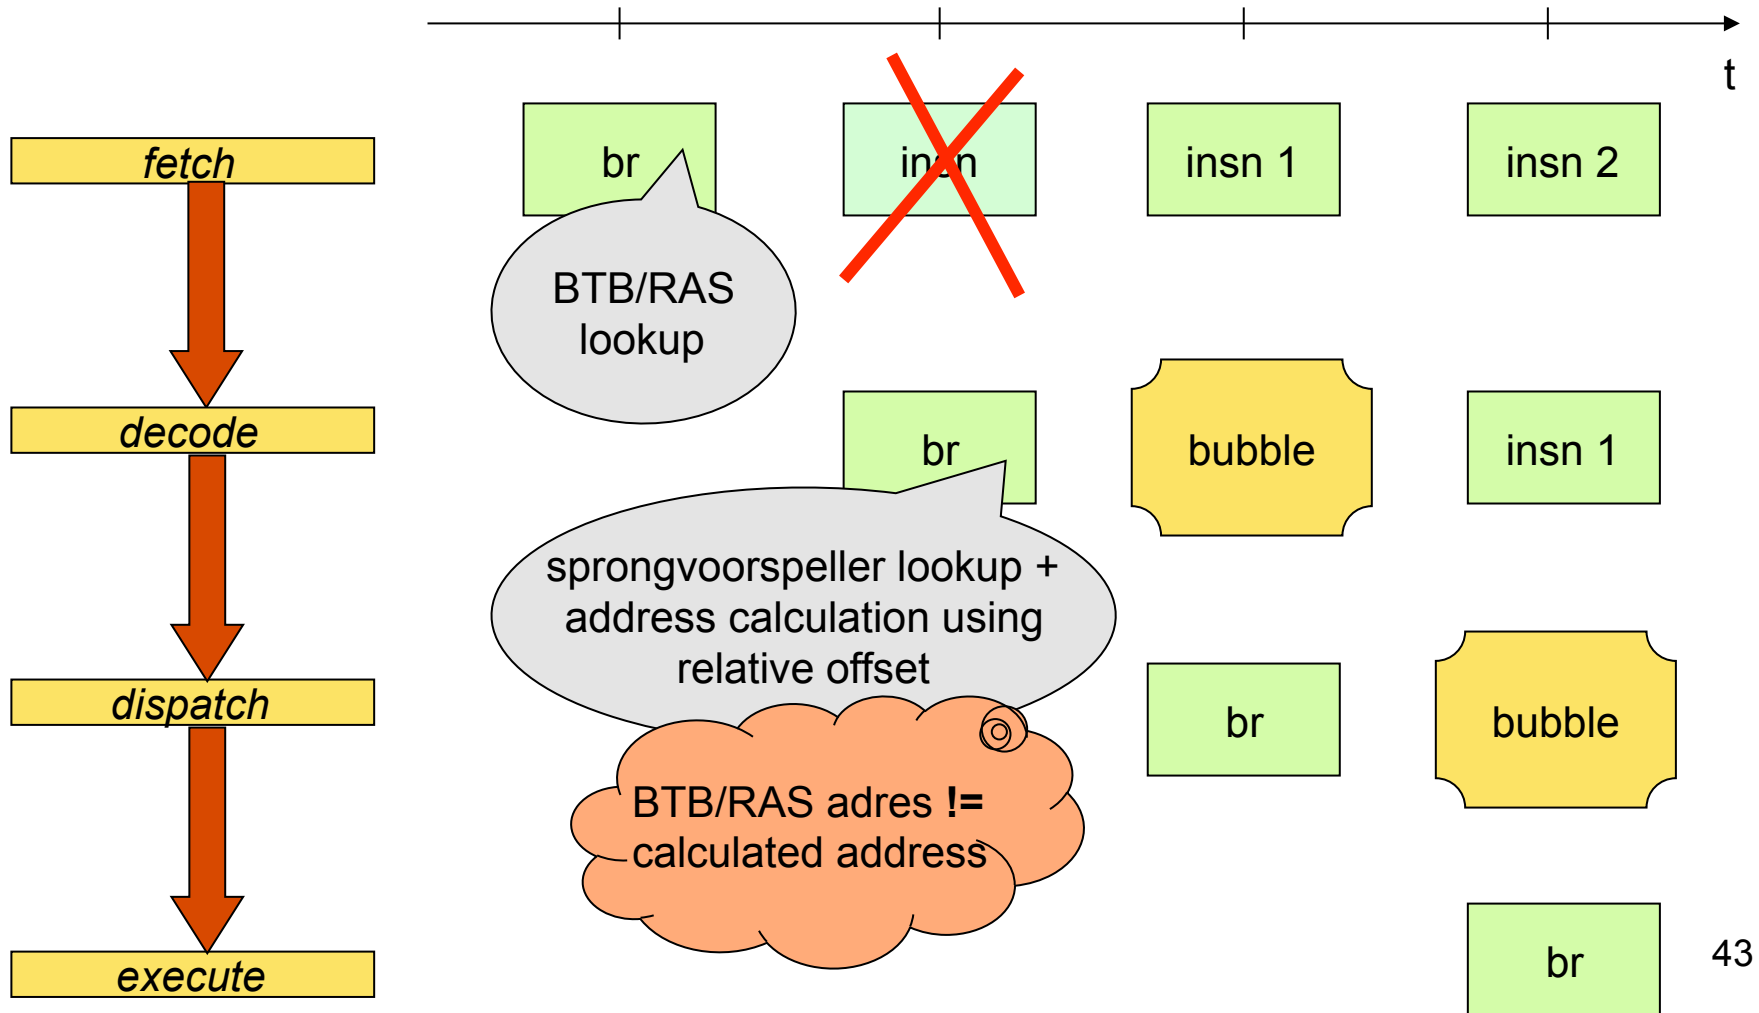
Branch prediction in the pipeline



Branch prediction predicts taken and target address is correct



Branch predictor predicts taken but target address is incorrect




Overview

- Branch direction prediction
 - Static
 - Dynamic
- Branch target prediction
- Branch misprediction recovery

Speculative execution

- Predict branch target address and start fetching and executing instructions along the predicted path
 - no completion of speculative insns!
- There might be multiple branches in flight along the predicted path
 - tags are added


Speculative execution




Pipeline squashing

- Branch outcome is known at execution
- If correctly predicted
 - Corresponding tags are deallocated
 - Instructions become non-speculative
- If mispredicted
 - Wrong-path instructions are nullified
 - Start fetching instructions along correct path

Speculative execution



Squash wrong-path insns and fetch insns along correct path



Literature

“Processor Microarchitecture: An Implementation Perspective” — A. González, F. Latorre, and G. Magklis

Morgan & Claypool Publishers,

Synthesis Lectures on Computer Architecture, 2010

“An Analysis of Correlation and Predictability: What Makes Two-Level Branch Predictors Work” —

M. Evers, S. J. Patel, R. S. Chappell, and Y. N. Patt

International Symposium on Computer Architecture (ISCA), June 1998